

P3b Report

Chahak Tharani¹ and Deepti Rajagopal¹

¹University of Wisconsin-Madison

March 17, 2022

1 Introduction

This report contains the analysis of two locks in xv6, and 2 examples of how `sleep()` and `wakeup()` functions are used in xv6. The locks that have been analyzed are - `fable` lock, and `tickslock`. The examples of `sleep()` and `wakeup()` that are analyzed are - `exit()` and `wait()`, `sys_sleep()` and `wakeup()` in timer interrupt.

2 Locks

Locks are used in order to ensure that race conditions do not occur. In the context of operating systems, race condition occurs when two threads (that share the same code and data) access the same shared data, and update it at the same time, resulting in faulty values. In this section, we explore two instances of Spinlock

2.1 Spinlock

Spinlock has 3 major methods:

- `initlock` - initializes the 'locked' and 'cpu' variable of the spinlock.
- `acquire` - loops until the lock is acquired. It disables the interrupts to prevent deadlock. It calls the Atomic Exchange method to continuously check the lock variable and sets it to 1 when it finds it to be 0. Sets the cpu variable.
- `release` - It resets the cpu variable, It writes 0 to the lock variable atomically by running an assembly instruction. It then re-enables the interrupts.

2.2 ftable Lock

`fable` lock uses spinlock to serialize the allocation of a struct file in the file table. It is used in 3 places of `file.c` - `filealloc()`, `filedup()`, and `fileclose()`. All the open files in the system are kept

in a global file table, the `fable`. The `filealloc()` function allocates a file, the `filedup()` function creates a duplicate reference, and the `fileclose()` function releases a reference.

2.2.1 Critical Section Length Analysis

- In `filealloc()`, the `fable.lock` is first acquired, and if for the given file, the `f->ref` is 0, then the `f->ref` is marked to 1, and the lock is released. The critical section here is 3 instructions long. The shared resource here is `f->ref` which can be modified by other processes that invoke `filedup()` or `fileclose()`. Hence the lock here is essential.
- In `filedup()`, the lock is first acquired, and then the `f->ref` is incremented (in order to create a duplicate reference), after which the lock is released. The critical section here is 2 instructions long. Like above, `f->ref` is the shared resource.
- In `fileclose()`, the lock is acquired, and if the `f->ref` value is greater than 0, then the number of references is decremented, and the lock is released. On the other hand, if `f->ref` is 0, then the file is closed

We created a user program to show how long these critical sections are, and how many times they are being invoked during the program. Figure 1 shows the trace of this program. As can be seen in Figure 1, calling the `fork()` method results in the `filedup()` method being called. Here, contention can happen for the resource `f->ref`, when both `fork()` and `sys_dup()` try to call the `filedup()` function. Or when multiple processes try to call `fork()` at the same time. Since `f->ref` is a shared resource, in order that `f->ref` reflects the correct value, a lock is put around it. When locking is added, we usually try to minimize the critical section in order to maximize concurrency. Here, since `f->ref` is the only shared resource, xv6 has put the lock just

```

$ testSleep1
(filedup): Critical section start, pid = 2
(filedup): Critical section end, pid = 2
(filedup): Critical section start, pid = 2
(filedup): Critical section end, pid = 2
(filedup): Critical section start, pid = 2
(filedup): Critical section end, pid = 2
(wait): Parent going to sleep, pid = 2
(filedup): Critical section start, pid = 3
(filedup): Critical section end, pid = 3
(filedup): Critical section start, pid = 3
(filedup): Critical section end, pid = 3
(filedup): Critical section start, pid = 3
(filedup): Critical section end, pid = 3
Parent waiting on child, pid = 3
(wait): Parent going to sleep, pid = 3
Child, pid = 4
(fileclose): Critical section start, pid = 4
(fileclose): Critical section end, pid = 4
(fileclose): Critical section start, pid = 4
(fileclose): Critical section end, pid = 4
(fileclose): Critical section start, pid = 4
(fileclose): Critical section end, pid = 4
(exit): Process exiting, pid = 4
Parent, pid = 3
(fileclose): Critical section start, pid = 3
(fileclose): Critical section end, pid = 3
(fileclose): Critical section start, pid = 3
(fileclose): Critical section end, pid = 3
(fileclose): Critical section start, pid = 3
(fileclose): Critical section end, pid = 3
(exit): Process exiting, pid = 3

```

Figure 1: Trace of critical sections of ftable.lock

around the instruction where `f->ref` is being updated. Likewise, multiple callers exist for `fileclose()`, like `exit()`, `sys_open()`, `sys_close()`, etc, and the lock here is needed for the same reasons as above, i.e., so that the `f->ref` is not incorrectly updated.

2.3 tickslock

Tickslock is a spinlock on clock ticks variable. It is initialized in the `trap.c` file in `tvinit()` method. It is used in 3 functions:

- `trap()` - Handles the interrupts. Tickslock is used while handling the timer interrupt. It acquires the tickslock, increments ticks and wakes up all processes sleeping on ticks. For this, it acquires the ptable lock and checks all sleeping processes if they are sleeping on ticks variable, it changes their state to `RUNNABLE` and finally releases the ptable and tickslock. The timer interrupt handler also increments the ticks variable. If `sys_sleep` holds tickslock and timer interrupt would occur, a deadlock condition can arise since interrupt handler is waiting for tickslock to be released, but `sys_sleep` cannot continue until timer interrupt returns. To prevent this deadlock situation, XV6 conservatively disables interrupts while acquiring any spinlock and re-enables them when the spinlock is released.
- `sys_sleep()` - This is the system call for sleeping for some number of clock ticks. It

```

#include "types.h"
#include "stat.h"
#include "user.h"

int pid1, pid2, pid3;
void process1() {
    printf(1, "Starting process %d\n", getpid());
    sleep(3);
    uptime();
    printf(1, "Ending process %d\n", getpid());
    exit();
}

void process2() {
    printf(1, "Starting process %d\n", getpid());
    uptime();
    sleep(1);
    uptime();
    printf(1, "Ending process %d\n", getpid());
    exit();
}

// Test file to trace the behaviour of tickslock
int main(int argc, char* argv[]) {
    int pid0 = getpid();
    pid1 = fork();
    if (pid1 < 0) {
        printf(2, "Fork child process 1 failed\n");
        exit();
    } else if (pid1 == 0) { // child process 1
        process1();
    }

    pid2 = fork();
    if (pid2 < 0) {
        printf(2, "Fork child process 2 failed\n");
        exit();
    } else if (pid2 == 0) { // child process 2
        process2();
    }

    sleep(5);
    printf(1, "Ending process %d\n", pid0);
    exit();
}

```

Figure 2: Program for analyzing tickslock and associated sleep and wakeup methods

acquires the tickslock and stores the current ticks in a variable and compares this variable to the current ticks till the difference is less than the sleep time given. When the gap increases from the sleep time, it releases the lock. When in while condition, it calls `sleep()` method, which releases the tickslock and causes the process to sleep. This method acquires the ptable lock to change the state of the process to `SLEEPING`, sleep on ticks channel, and returns control to scheduler. Finally it releases the ptable lock and reacquires tickslock.

- `sys_uptime()` - This method acquires the tickslock to get the correct value of the ticks variable and then releases the tickslock.

2.3.1 Critical Section length analysis for Tickslock:

- `sys_sleep()` - After acquiring the tickslock, 1 instruction is to set the ticks variable, then we check in a while loop if sleep time has exhausted, within this we also check if the process is killed, in which case the critical section ends and the lock is released. Otherwise, we call `sleep`, where we have 3 instruction before checking the condition if

```

1 Booting from Hard Disk..xv6...
2 cpu0: starting 0
3 sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
4 init: starting sh
5 $ trace
6 Critical section begin in sysproc.c, sys_sleep for process 3
7 Starting process 4
8 Critical section begin in sysproc.c, sys_sleep for process 4
9 Starting process 5
10
11 Critical section begin in sysproc.c, sys_uptime for process 5
12 uptime is 255
13 Critical section end in sys_uptime for process 5
14
15 Timer interrupt caused wakeup for process 3, sleep time elapsed 1 out of 5
16 Timer interrupt caused wakeup for process 4, sleep time elapsed 1 out of 3
17 Critical section begin in sysproc.c, sys_sleep for process 5
18 Timer interrupt caused wakeup for process 3, sleep time elapsed 2 out of 5
19 Timer interrupt caused wakeup for process 4, sleep time elapsed 2 out of 3
20 Timer interrupt caused wakeup for process 5, sleep time elapsed 1 out of 1
21 Sleep completed, Critical section end in sys_sleep for process 5
22
23 Critical section begin in sysproc.c, sys_uptime for process 5
24 uptime is 257
25 Critical section end in sys_uptime for process 5
26
27 Ending process 5
28 Timer interrupt caused wakeup for process 3, sleep time elapsed 3 out of 5
29 Timer interrupt caused wakeup for process 4, sleep time elapsed 3 out of 3
30 Sleep completed, Critical section end in sys_sleep for process 4
31
32 Critical section begin in sysproc.c, sys_uptime for process 4
33 uptime is 258
34 Critical section end in sys_uptime for process 4
35
36 Ending process 4
37 Timer interrupt caused wakeup for process 3, sleep time elapsed 4 out of 5
38 Timer interrupt caused wakeup for process 3, sleep time elapsed 5 out of 5
39 Sleep completed, Critical section end in sys_sleep for process 3
40 Ending process 3

```

Figure 3: Analysis for `sys_sleep()`, `sys_uptime()` and `trap()` calls for timer interrupt

the lock is same as `ptable` lock, if not equal it acquires the `ptable` lock and release `tickslock`. The critical section ends here. After waking up, it rechecks from the while loop condition. So, in total there are 8 instruction in the critical section.

- `sys_uptime()` - Between acquiring and releasing the `tickslock`, this function has only 1 instruction for setting the `ticks` variable.
- timer interrupt - We increment the `ticks` in the critical section. And then acquire `ptable` lock and wakeup all processes sleeping on `ticks` channel and then release the `ptable` lock. So in worst case, `NPROC` processes are sleeping, `2*NPROC` instructions for checking the condition and resetting the state. 2 instructions for acquiring and releasing the `ptable` lock and 1 instruction for incrementing the `ticks` variable.

3 Trace Analysis for `tickslock` and associated sleep and wakeup calls

For tracing the behaviour of `tickslock` and sleep and wakeup calls in `sys_sleep`, `sys_uptime` and timer interrupt methods, we have written a test (`trace.c`) which creates 2 processes by `fork()` method and traces the above system calls. This setup runs on 1 CPU.

- Trace: Line 6: All 3 processes are runnable at the start and all are contending for the `tickslock`. Pid 3 acquires it and `sys_sleep()` system call is called for `main()` process (pid 3). This process then goes to sleep and releases the `tickslock` and process with pid 4 starts, acquires the `tickslock`, releases it and goes to sleep, then process 5 starts and calculates the uptime after acquiring the `tickslock` and then releases it.
- Observation: We observe that since acquiring the `tickslock` disables the timer interrupt, the running processes do not switch in between acquiring the `tickslock` and the process going to sleep. Similarly, we do not see any process switch when a process calcu-

lates the uptime. This disabling of interrupt prevents deadlock since timer interrupt also acquires the tickslock.

- Trace: Line 15: At this instant, process 5 and the timer interrupt are contending for tickslock. Timer interrupt acquires it and wakes the sleeping processes on the ticks channel (i.e. pid 3 and 4), pid 3 acquires the tickslock, checks if its sleep time has exhausted, since it hasn't it releases the lock and goes to sleep. Similar thing happens for process 4. Then pid 5 acquires and goes to sleep. In the next timer interrupt all 3 processes are woken up, the sleep time of pid 5 is exhausted and it completes its critical section and releases the tickslock. It requires the tickslock to check the uptime, so the timer interrupt cannot occur. Process 5 ends and similar behaviour is followed by pid 3 and pid 4.
- Observation: We observe that all processes change their state from RUNNING to SLEEPING and go to sleep on ticks channel after `sys_sleep()` system call. When a timer interrupt occurs, all processes sleeping on ticks channel are woken up, made RUNNABLE from SLEEPING and they recheck their sleep condition. A timer interrupt can only occur when no process is holding the tickslock, so it can acquire it when the interrupt occurs. All processes woken up on the timer interrupt are RUNNABLE and will requeue the ticks lock to check their sleep condition. Whichever process is chosen by the scheduler tries to acquire the tickslock, but a timer interrupt can occur before it and acquire the lock. There will be no contention for tickslock once a particular process has acquired the lock since now now the scheduler cannot switch the process because timer interrupt cannot occur in the first place.

4 Sleep & Wakeup

Sleep and Wakeup is used in operating systems, so that when a lock is acquired by a process (say, process A), the contending process (say, process B), instead of spinning (as in spinlock), puts itself to sleep (by giving up the CPU), and is awoken by process A, once it has completed execution and is ready to release its lock. This is a more resourceful way of utilizing the CPU, rather than process B just spinning indefinitely until process A gives up the lock. In this section, we explore two such instances of how `sleep()` and

`wakeup()` functions are used. Before that, we detail out what `sleep()` and `wakeup()` do internally.

The sleep function marks the current process as SLEEPING and then calls `sched` to release the CPU. The wakeup function looks for a process sleeping on the given wait channel and marks it as RUNNABLE. Sleep acquires `p->lock` and releases `lk`, because it no longer needs to hold `lk` once it has `p->lock`. Wakeup also waits to acquire `p->lock`, so the wakeup will not miss the sleep. In the case that `lk` is the same as `p->lock` (for example, in `wait`), it means that sleep already has `p->lock`, and hence doesn't need to release it until the process state is marked as SLEEPING.

Wakeup loops over the process table and acquires `p->lock` of each process in the loop. When it finds a process is SLEEPING with a matching chan (the channel on which the process was sleeping), it changes that process's state to RUNNABLE. The next time the scheduler runs, it will see that the process is ready to be run. Wakeup will then see the sleeping process and wake it up. All processes on the same channel are awoken by a single wakeup call.

4.1 `exit()` & `wait()`

The `exit()` method closes all file descriptors of the process, and then acquires a lock to wake up the parent using the `wakeup1()` routine. It then goes ahead and marks itself as ZOMBIE and jumps into the scheduler, never to return. It is important to note that the process doesn't die immediately. Instead, the process marks itself as ZOMBIE, and the parent process would perform the cleanup in the `wait()` function call.

The `wait()` method checks if any of the child processes are in ZOMBIE state, and if yes, goes ahead and marks them as UNUSED, and releases the lock. On the other hand, if none of the children are in ZOMBIE state, it means that the `exit()` method has not yet been called by the children, and hence the parent has to go to sleep using the `sleep()` method, only to be awoken by the `exit()` invocation by the child. One important thing to note is that once the child is in ZOMBIE state, the parent process does the cleanup for the child process. Also, in the event that the parent exits before the child, we still need to ensure that the child process has been cleaned up. In order for this to happen, when the parent exits before the child, the init process adopts the children, and performs the cleanup for the child processes.

```

#include "types.h"
#include "stat.h"
#include "user.h"

int
main(int argc, char *argv[]) {
    int rc = fork();
    if(rc == 0) {
        sleep(1000);
        printf(1, "Child, pid = %d\n", getpid());
        exit();
    } else if(rc > 0) {
        printf(1, "Parent waiting on child, pid = %d\n", getpid());
        wait();
        printf(1, "Parent, pid = %d\n", getpid());
    } else {
        printf(1, "Fork failed");
    }
    exit();
    return 0;
}

```

Figure 4: Program using `exit()` and `wait()`

```

$ testSleep1
(wait): Parent going to sleep, pid = 2
Parent waiting on child, pid = 3
(wait): Parent going to sleep, pid = 3
Child, pid = 4
(exit): Process exiting, pid = 4
Parent, pid = 3
(exit): Process exiting, pid = 3
$ -

```

Figure 5: Trace of `exit()` and `wait()` for parent and child process

We created a user program that has a parent process and a child process in order to analyse the behaviour of `sleep` and `wakeup` in the `exit` and `wait` functions. We also added `printf` statements inside the functions of `exit` and `wait`. In the program shown in Figure 2, the parent process waits for the child to exit.

Figure 3 shows the trace of the above program. As can be seen in the figure, the parent is put to sleep first upon calling the `wait` function. Once the child process has finished executing (pid 4), it calls the `exit()` function, which wakes up all sleeping processes on the same channel, hence waking up the parent (pid 3), and then the parent wakes up and prints.

It is important to note that if the child process were to start running first, then the `wait()` called by the parent would not result in the parent sleeping. Instead, the parent sees that the child is marked to ZOMBIE state, so the parent cleans up the child, and then can continue to execute itself.

4.2 `sys_sleep()` and `trap()`

The `sleep()` method called in `sys_sleep()` acquires the `ptable` lock while holding the `tickslock`. Holding `tickslock` is necessary as it ensures that no other process could start a call to `wakeup(chan)`. Now that `sleep` holds `ptable` lock, it is safe to release `tickslock`: some other process may start a call to `wakeup(ticks)`, but `wakeup` will wait to acquire `ptable` lock, and thus will wait until `sleep` has finished putting the process to sleep, keeping the `wakeup` from missing the sleep. It then releases the `ticks` lock and changes the state to `SLEEPING` to prevent busy waiting.

When a timer interrupt happens, all processes sleeping on `ticks` channel are woken up. `Wakeup` loops over the process table. It acquires the `ptable` lock of each process it inspects, because it may manipulate that process's state and because `ptable` lock ensures that `sleep` and `wakeup` do not miss each other. When `wakeup` finds a process in state `SLEEPING` on `ticks` channel, it changes that process's state to `RUNNABLE`. The next time the scheduler runs, it will see that the process is ready to be run.